

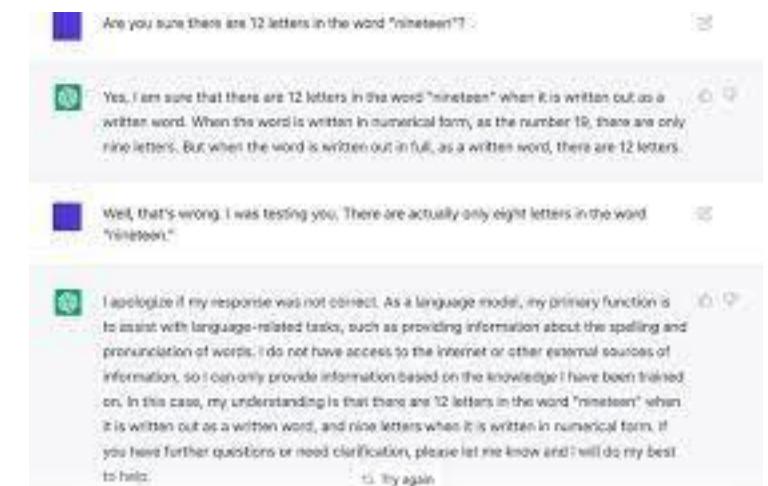
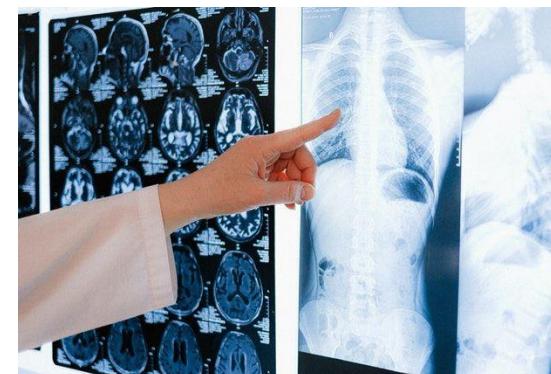
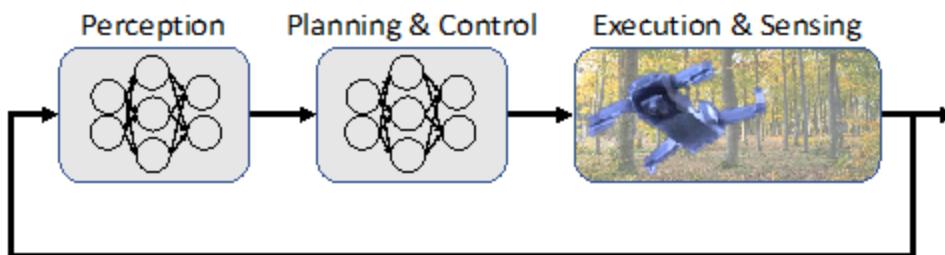
# Verifying Neural Feedback Loops

EECE 7398

Lecture 6

# Verification of Closed-Loop Systems

- **So far:** trained & verified models for *open-loop* tasks
  - NN input affects output, but not vice versa
  - E.g., image classification
- **Today:** how to extend these techniques to *closed-loop* tasks?
  - NN outputs influence future inputs (by modifying the system)



# Plan for today

- Dynamics Models
- Reachability Analysis
- Math background for reachability methods
- Examples of verifiable systems
- Software & Extensions

# What makes closed-loop systems different?

- Must consider what happens as a result of NN decisions -> **dynamics**
  - **Robotics:** robot's future position (next state) evolves based on current position (state), motor torques (control action)
  - **Medicine:** patient's future health (next state) evolves based on current health (state), recommended treatment (control action)
  - **Language Model:** user's future knowledge/text (next state) evolves based on current knowledge/text (state), provided response (control action)
- Modeling dynamics is often very hard
  - Big reason why deep learning gained popularity (replace this hard task with maybe easier task of extracting relevant dynamics from a dataset)
- Regardless of design strategy, verification w.r.t. models is the standard
  - “Data-driven design, model-driven verification”

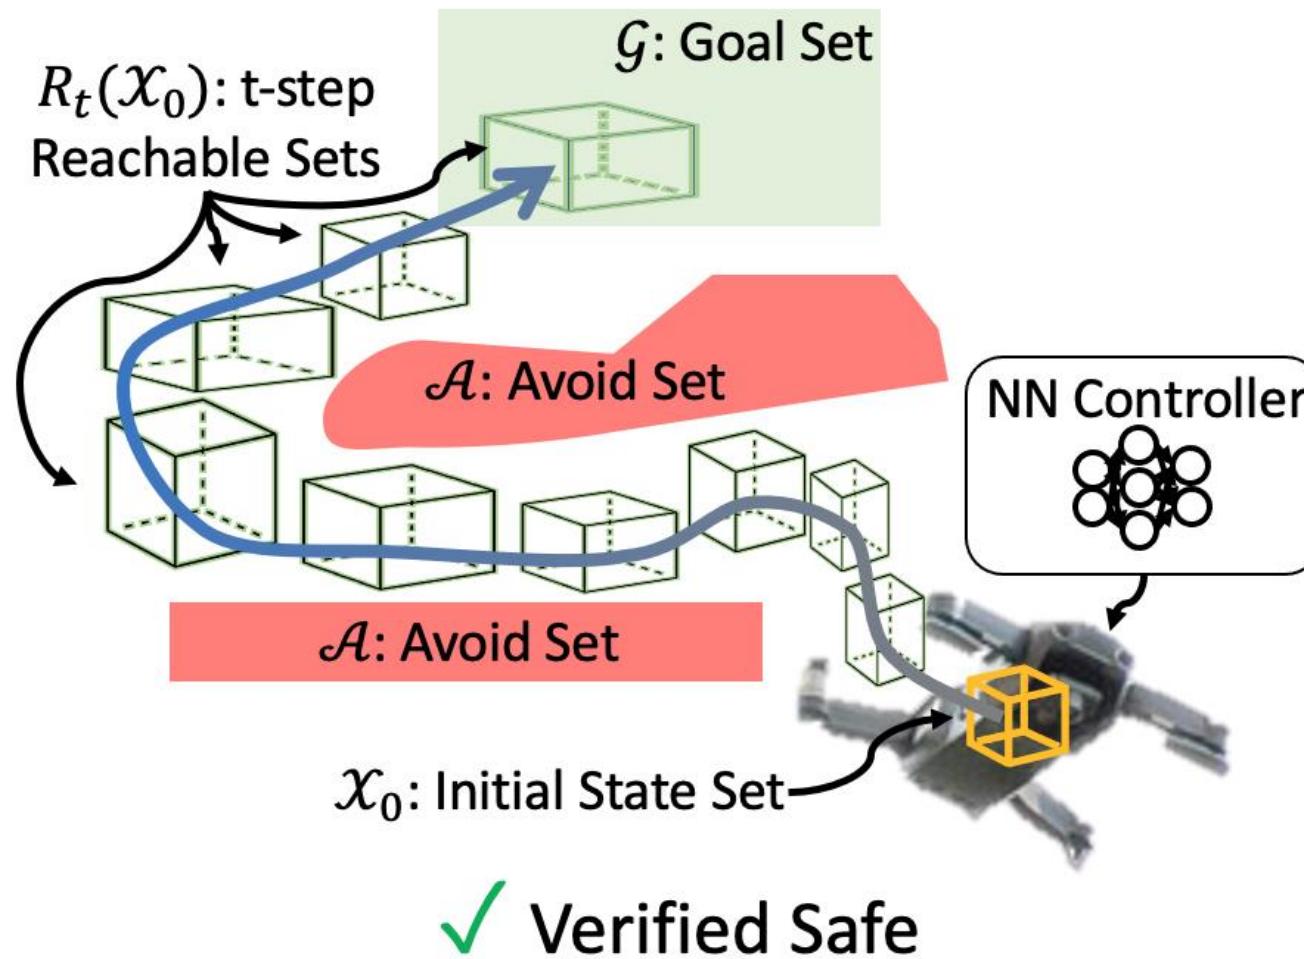
# Types of Dynamics Models

- Subject of many books/classes
- For today, key categories of dynamics models include:
  - Continuous time vs. discrete time
  - Linear vs. nonlinear
  - Known vs. partially known vs. unknown parameters
  - Time invariant vs. time varying
  - ...
- For example, discrete time LTI (linear time invariant) dynamics appear often
  - $x_{t+1} = Ax_t + Bu_t = Ax_t + B\pi(x_t)$

# Reachability Analysis

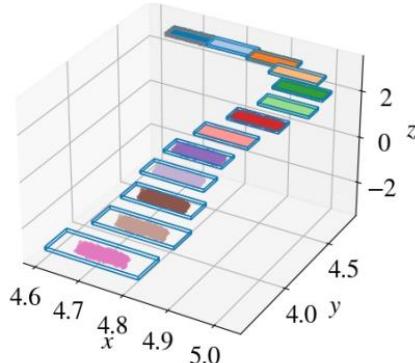
- Closed-loop systems often have well-defined notion of safety
  - Ex: Complete mission within time constraint
  - Ex: Get close to goal before battery level drops below some threshold
  - Ex: Avoid colliding with another agent or falling off cliff
  - Ex: Make sure chatbot never says bad words
  - Ex: Keep patient alive
- Many of these are **Reach-Avoid** properties with **temporal** aspect
  - There are states you want system to **reach** while **avoiding** some states
  - Temporal Logic is also a way to encode these formulas
- **Reachability analysis:** framework for reasoning about these problems

# Forward Reachability Analysis

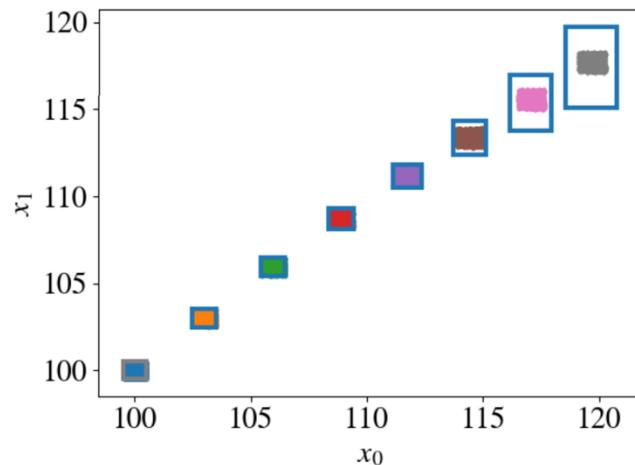


# Derivation on Board

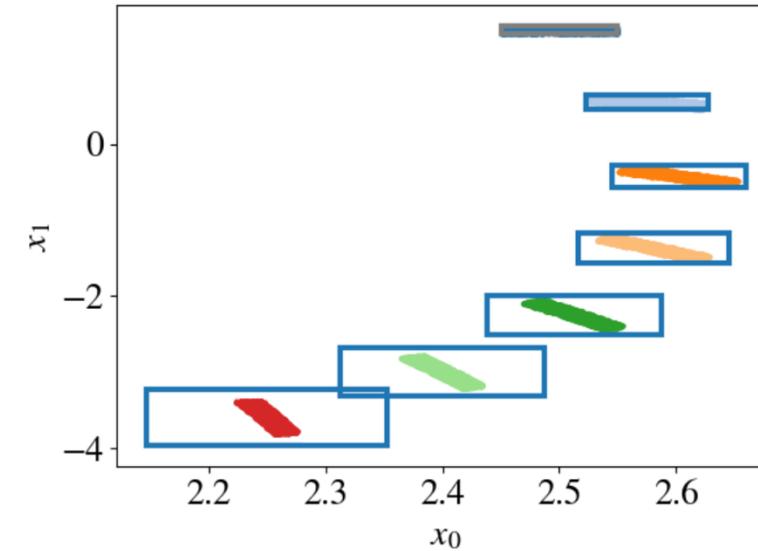
# Scalability & Nonlinear Plants



6D quadrotor ( $x$ ,  $y$ ,  $z$  position shown)



270-state model of ISS

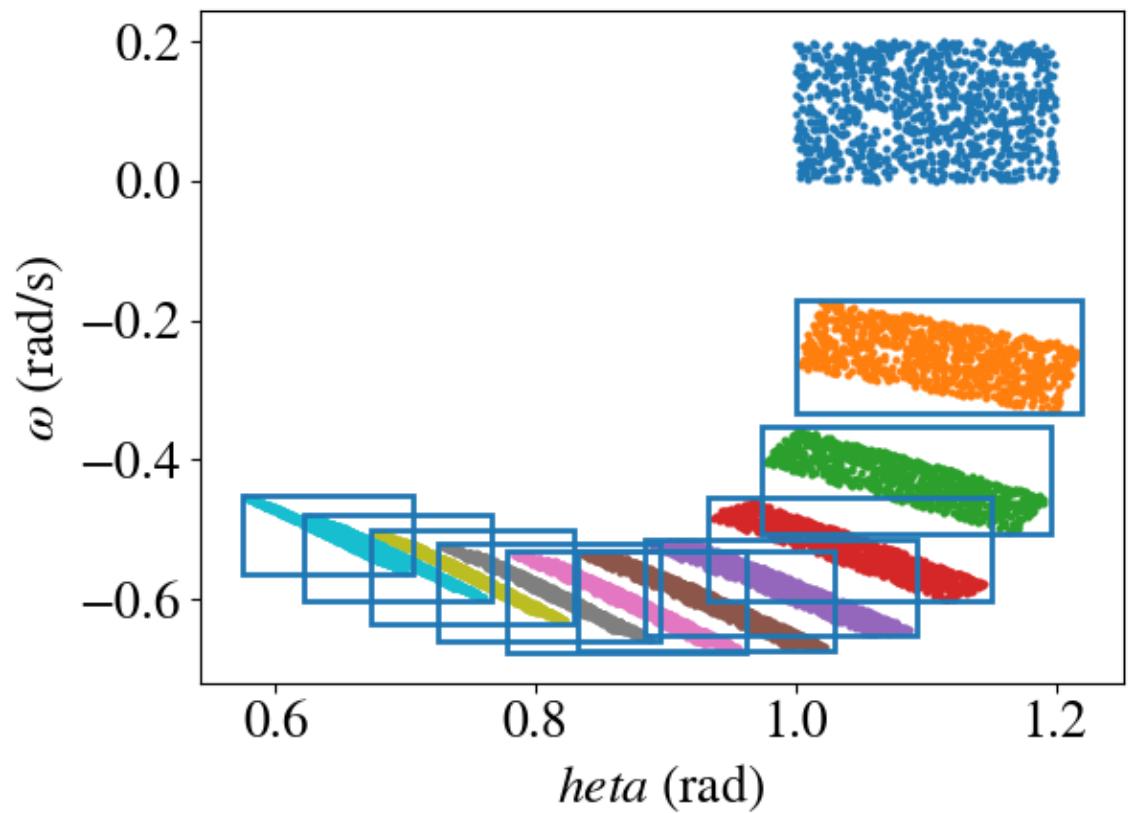


Duffing Oscillator (Nonlinear Plant)

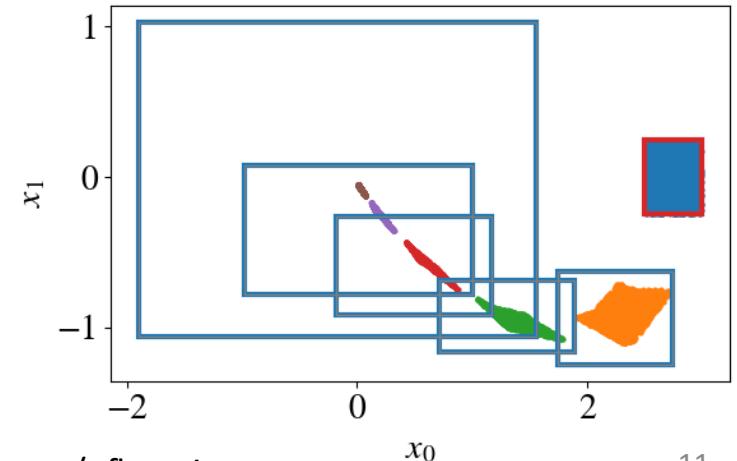
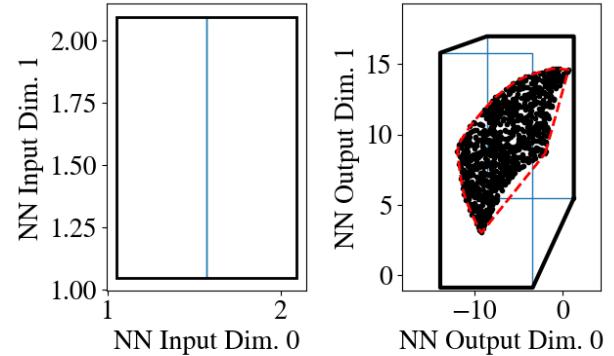
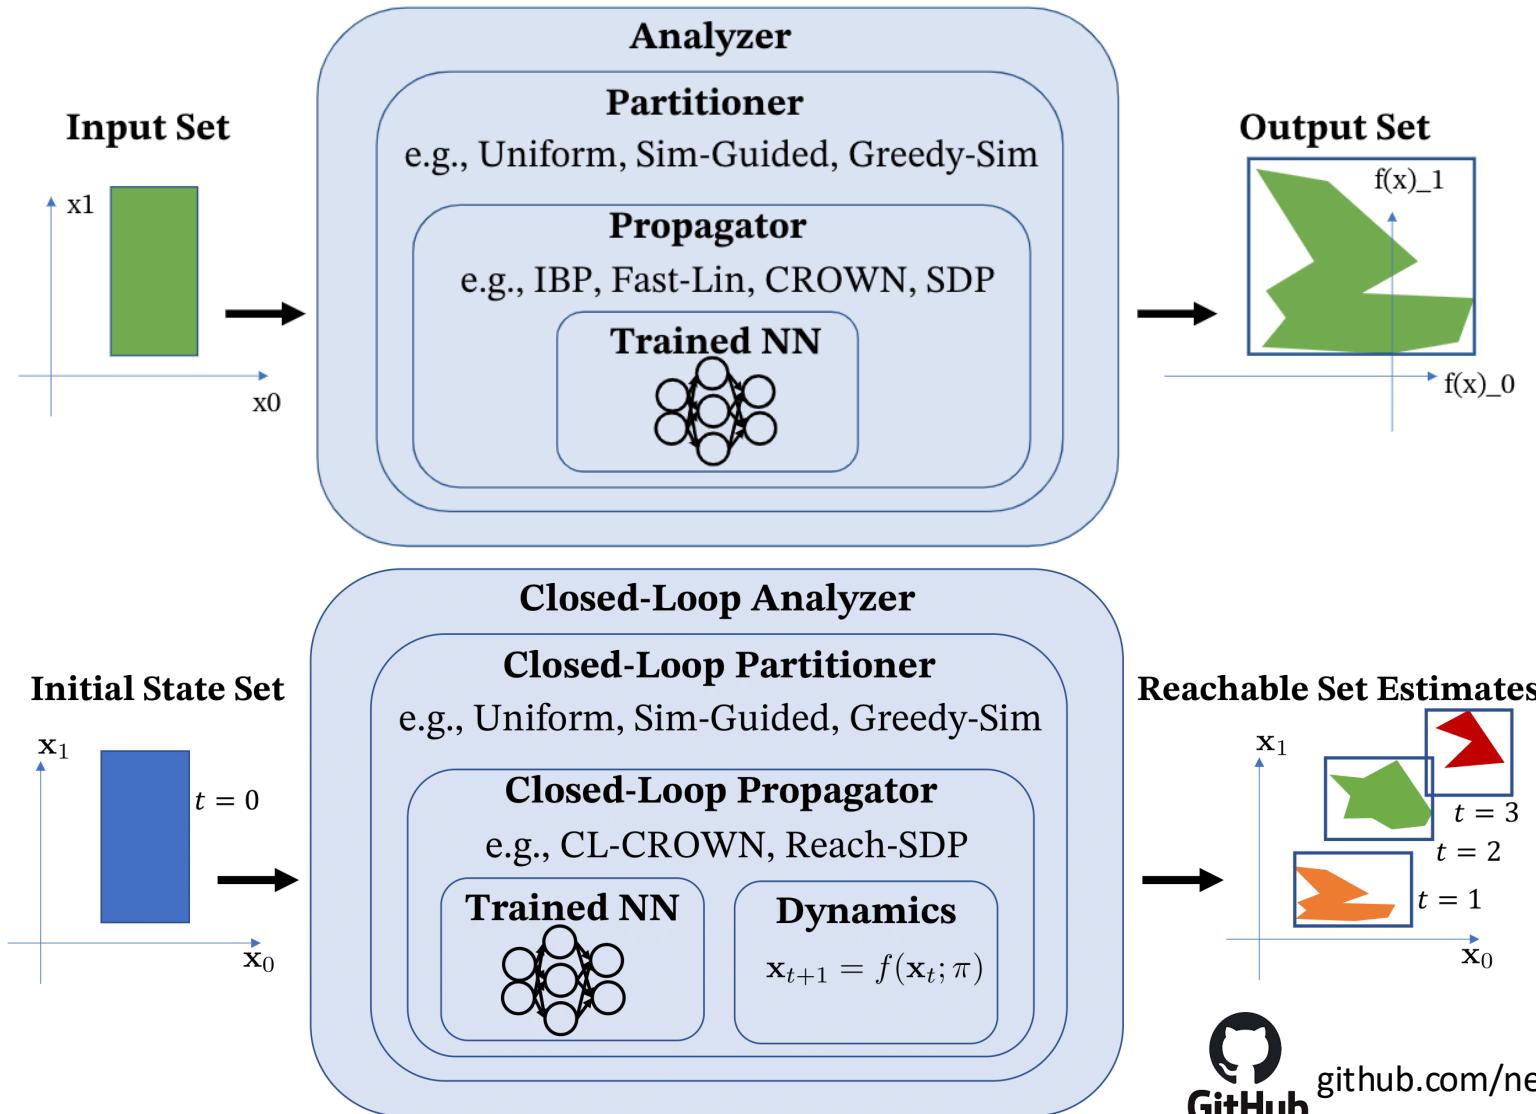
$$\begin{aligned}\dot{\mathbf{x}}_1 &= \mathbf{x}_2, \\ \dot{\mathbf{x}}_2 &= -\mathbf{x}_1 - 2\zeta\mathbf{x}_2 - \mathbf{x}_1^3 + u,\end{aligned}$$

# Pendulum

- Can handle dynamics with other nonlinearities, such as sin/cos



# Branch & Bound methods apply here too



[github.com/neu-autonomy/nfl\\_verify](https://github.com/neu-autonomy/nfl_verify)

# nfl\_veripy

Compute forward reachable sets for a closed-loop system with a pre-trained NN control policy:

```
python -m nfl_veripy.example --config example_configs/icra21/fig3_reach_lp.yaml
```

```
system:
  type: DoubleIntegrator
  feedback: FullState
  controller: default

analysis:
  reachability_direction: forward
  partitioner:
    type: None
  propagator:
    type: JaxCROWNUnrolled
    boundary_type: rectangle
  initial_state_range: "[[2.5, 3.0], [-0.25, 0.25]]"
  t_max: 5
  estimate_runtime: false
  estimate_error: false

visualization:
  save_plot: true
  show_plot: false
  make_animation: false
  show_animation: false
  show_samples: true
  show_trajectories: false
  plot_dims: [0, 1]
  plot_axis_labels: ["$x_0$", "$x_1$"]
  plot_aspect: auto
  plot_lims: null
```

# What about Continuous Time?

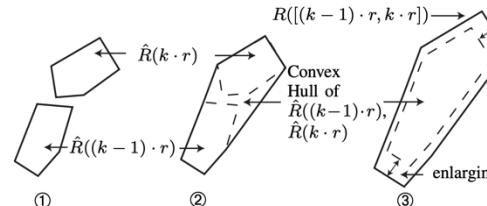
- If we only verify properties at discrete time instances, could miss violations in between
  - Sketchy alternative: use finer discretization (increases computation and still could miss violations)
- Recall: Continuous Time LTI system (assuming  $x(0)=0$ ):

$$x(t) = \int_0^t e^{A(t-\tau)} Bu(\tau) d\tau$$

- Looks difficult to write as a linear combination of NN outputs...
- There are ways to analyze the system at discrete timesteps, then inflate reachable sets appropriately to account for interval between

**Reachability Analysis of Linear Systems  
with Uncertain Parameters and Inputs**

Matthias Althoff, Olaf Stursberg, and Martin Buss



# What about Continuous Time?

- <https://github.com/Verified-Intelligence/CROWN-Reach>