

Forward Reachability Analysis of Neural Feedback Systems

EECE/CS 7268: Verifiable Machine Learning

Michael Everett

1 System Dynamics

Consider a system with:

$$\mathbf{x}_{t+1} = p(\mathbf{x}_t, \mathbf{u}_t) \quad (\text{open-loop dynamics}) \quad (1)$$

$$\mathbf{u}_t = \pi(\mathbf{x}_t) \quad (\text{NN controller}) \quad (2)$$

By substituting the controller equation into the dynamics, we obtain the closed-loop system:

$$\mathbf{x}_{t+1} = p(\mathbf{x}_t, \pi(\mathbf{x}_t)) = f(\mathbf{x}_t, \pi) \quad (\text{closed-loop dynamics}) \quad (3)$$

Figure 1 illustrates this system structure, showing how the neural network controller π and plant dynamics p interact to form the closed-loop dynamics f .




Figure 1: Closed-loop system with neural network controller

Example: $\mathbf{x}_{t+1} = A\mathbf{x}_t + B\pi(\mathbf{x}_t)$

2 Multi-Step Closed-Loop Dynamics

Given the closed-loop dynamics $\mathbf{x}_{t+1} = f(\mathbf{x}_t, \pi)$ and an initial state \mathbf{x}_t , we can compute future states by repeatedly applying f :

$$\mathbf{x}_{t+2} = f(\mathbf{x}_{t+1}, \pi) \quad (4)$$

$$= f(f(\mathbf{x}_t, \pi), \pi) \quad (5)$$

$$= f^2(\mathbf{x}_t, \pi) \quad (6)$$

This iterative application of the closed-loop dynamics over multiple timesteps is represented in Figure 2, where each f block corresponds to the dashed box from Figure 1. By chaining these blocks together, we can compute the state evolution over any time horizon T .




Figure 2: Multi-step closed-loop dynamics over T timesteps, where each f block represents the combined controller and plant dynamics

3 Reachable Sets

When the closed-loop dynamics f is deterministic and the initial state \mathbf{x}_t is known precisely, the future system behavior is perfectly described by a single trajectory. However, in practice, we often only know that the initial state lies within some set \mathcal{X}_t (e.g., due to estimation uncertainty or system variability). In this case, we need to characterize all possible future states the system could reach from any starting condition in \mathcal{X}_t . This leads to the concept of reachable sets.

Definition 1. *The reachable set at the initial time t is simply the initial set:*

$$\mathcal{R}_0(\mathcal{X}_t) = \mathcal{X}_t \quad (7)$$

Definition 2. *The 1-step reachable set from a given set \mathcal{X}_t is all states the system could be in at time $t + 1$, assuming it starts from some state in \mathcal{X}_t :*

$$\mathcal{R}_1(\mathcal{X}_t) = \{f(\mathbf{x}_t, \pi) \mid \forall \mathbf{x}_t \in \mathcal{X}_t\} \quad (8)$$

Definition 3. *The T -step reachable set contains all states reachable after exactly T steps from any initial state in \mathcal{X}_t :*

$$\mathcal{R}_T(\mathcal{X}_t) = \{f^T(\mathbf{x}_t, \pi) \mid \forall \mathbf{x}_t \in \mathcal{X}_t\} \quad (9)$$

We sometimes want to know the set of all states reachable within time horizon T , denoted as $\mathcal{R}_{0:T}(\mathcal{X}_t)$, which is the union of reachable sets from time $t = 0$ to $t = T$:

$$\mathcal{R}_{0:T}(\mathcal{X}_t) = \bigcup_{\tau=0}^T \mathcal{R}_\tau(\mathcal{X}_t) \quad (10)$$




Figure 3: Illustration of reachable sets evolving from left to right. \mathcal{X}_t is the initial set, \mathcal{R}_1 and \mathcal{R}_2 are connected non-convex sets, while \mathcal{R}_3 becomes disconnected.

4 Why Care About $\mathcal{R}_T(\mathcal{X}_t)$?

- If we have a goal region \mathcal{G} :

$$\mathcal{R}_T(\mathcal{X}_t) \subseteq \mathcal{G} \Rightarrow \text{system is guaranteed to reach the goal at timestep } T \quad (11)$$

- If we have an avoid region \mathcal{A} :

$$\mathcal{R}_{0:T}(\mathcal{X}_t) \cap \mathcal{A} = \emptyset \Rightarrow \text{system is guaranteed to not enter } \mathcal{A} \text{ any time in next } T \text{ steps} \quad (12)$$

5 Computing Reachable Sets

Computing \mathcal{R}_T exactly is intractable (NP-complete) [1].

Instead of computing exact reachable sets, we can calculate over-approximations or bounds, denoted as $\overline{\mathcal{R}}_T$, which are guaranteed to contain the true reachable set: $\mathcal{R}_T \subseteq \overline{\mathcal{R}}_T$. These bounds are valuable because they maintain soundness for verification:

- If $\overline{\mathcal{R}}_T(\mathcal{X}_t) \subseteq \mathcal{G}$, then we can guarantee $\mathcal{R}_T(\mathcal{X}_t) \subseteq \mathcal{G}$, meaning the true system is guaranteed to reach the goal.
- If $\overline{\mathcal{R}}_{0:T}(\mathcal{X}_t) \cap \mathcal{A} = \emptyset$, then we can guarantee $\mathcal{R}_{0:T}(\mathcal{X}_t) \cap \mathcal{A} = \emptyset$, meaning the true system is guaranteed to avoid the obstacle.

To compute such bounds, we formulate optimization problems that find the extremal values of the state in different directions:

$$\max_{\mathbf{x}_t \in \mathcal{X}_t} \mathbf{c}^T \mathbf{x}_{t+T} \quad \text{s.t.} \quad \mathbf{x}_{t+T} = f^T(\mathbf{x}_t, \pi) \quad (13)$$

These optimization problems can be solved using various techniques from neural network verification literature. This connection to neural network verification is natural because the closed-loop system f with a neural controller π forms a computational graph similar to a neural network. The multi-step dynamics f^T can be viewed as unrolling this computational graph T times, analogous to unrolling a recurrent neural network. Consequently, tools developed for neural network verification—such as MILP, LP, CROWN, IBP, and β -CROWN—are directly applicable to bounding reachable sets for neural feedback control systems.

6 Getting Hyper-Rectangle Bounds

To find a hyperrectangle bound on the reachable set \mathcal{R}_T , we can choose the objective vectors \mathbf{c} in our optimization problem to maximize or minimize each component of the state vector \mathbf{x}_{t+T} . This allows us to compute the tightest hyperrectangle (axis-aligned box) that contains the reachable set.

For example, to get the upper bound on the first component of the state, $\bar{x}_{t+T,1}$, we can let:

$$\mathbf{c}^T = [1 \ 0 \ \cdots \ 0] \quad (14)$$

In other words, we are solving the following optimization problem:

$$\bar{x}_{t+T,1} = \max_{\mathbf{x}_t \in \mathcal{X}_t} x_{t+T,1} \quad \text{s.t.} \quad \mathbf{x}_{t+T} = f^T(\mathbf{x}_t, \pi) \quad (15)$$

This same logic can be repeated for each element of the state to get the upper bound vector $\bar{\mathbf{x}}_{t+T}$:

$$\mathbf{c}^T = [1 \ 0 \ \cdots \ 0] \quad \text{to get} \quad \bar{x}_{t+T,1} = \max_{\mathbf{x}_t \in \mathcal{X}_t} x_{t+T,1} \quad (16)$$

$$\mathbf{c}^T = [0 \ 1 \ \cdots \ 0] \quad \text{to get} \quad \bar{x}_{t+T,2} = \max_{\mathbf{x}_t \in \mathcal{X}_t} x_{t+T,2} \quad (17)$$

$$\vdots \quad (18)$$

$$\mathbf{c}^T = [0 \ 0 \ \cdots \ 1] \quad \text{to get} \quad \bar{x}_{t+T,n} = \max_{\mathbf{x}_t \in \mathcal{X}_t} x_{t+T,n} \quad (19)$$

The upper bound vector is then constructed as:

$$\bar{\mathbf{x}}_{t+T} = [\bar{x}_{t+T,1}, \bar{x}_{t+T,2}, \dots, \bar{x}_{t+T,n}]^T \quad (20)$$

Then, we can flip the sign of the objective vector to get the lower bound vector $\underline{\mathbf{x}}_{t+T}$. For instance, to get the lower bound on the first component, $\underline{x}_{t+T,1}$, we use $\mathbf{c}^T = [-1 \ 0 \ \dots \ 0]$ in our maximization framework:

$$\max_{\mathbf{x}_t \in \mathcal{X}_t} \mathbf{c}^T \mathbf{x}_{t+T} = \max_{\mathbf{x}_t \in \mathcal{X}_t} (-1) \cdot x_{t+T,1} = - \min_{\mathbf{x}_t \in \mathcal{X}_t} x_{t+T,1} \quad (21)$$

Thus, $\underline{x}_{t+T,1} = \min_{\mathbf{x}_t \in \mathcal{X}_t} x_{t+T,1}$. Repeating for all components:

$$\mathbf{c}^T = [-1 \ 0 \ \dots \ 0] \quad \text{to get} \quad -\underline{x}_{t+T,1} = \max_{\mathbf{x}_t \in \mathcal{X}_t} (-x_{t+T,1}) \Rightarrow \underline{x}_{t+T,1} = \min_{\mathbf{x}_t \in \mathcal{X}_t} x_{t+T,1} \quad (22)$$

$$\mathbf{c}^T = [0 \ -1 \ \dots \ 0] \quad \text{to get} \quad -\underline{x}_{t+T,2} = \max_{\mathbf{x}_t \in \mathcal{X}_t} (-x_{t+T,2}) \Rightarrow \underline{x}_{t+T,2} = \min_{\mathbf{x}_t \in \mathcal{X}_t} x_{t+T,2} \quad (23)$$

$$\vdots \quad (24)$$

$$\mathbf{c}^T = [0 \ 0 \ \dots \ -1] \quad \text{to get} \quad -\underline{x}_{t+T,n} = \max_{\mathbf{x}_t \in \mathcal{X}_t} (-x_{t+T,n}) \Rightarrow \underline{x}_{t+T,n} = \min_{\mathbf{x}_t \in \mathcal{X}_t} x_{t+T,n} \quad (25)$$

The lower bound vector is constructed as:

$$\underline{\mathbf{x}}_{t+T} = [x_{t+T,1}, x_{t+T,2}, \dots, x_{t+T,n}]^T \quad (26)$$

Once we have computed these bounds, we have a hyperrectangle $[\underline{\mathbf{x}}_{t+T}, \bar{\mathbf{x}}_{t+T}]$ that is guaranteed to contain the true reachable set $\mathcal{R}_T(\mathcal{X}_t)$.




Figure 4: Illustration of hyperrectangle bounds (dashed red) for a 1-step reachable set. The true reachable set $\mathcal{R}_1(\mathcal{X}_t)$ (blue) has a non-convex shape and is fully contained within the hyperrectangle bounds, which form a rectangular over-approximation. The bounds are defined by the extreme points $\underline{\mathbf{x}}_{t+1}$ and $\bar{\mathbf{x}}_{t+1}$ (red dots), which are computed by solving optimization problems for the minimum and maximum values in each dimension.

7 Bounds in Matrix Form

Having computed the hyperrectangle bounds in Section 6, we can express these bounds more compactly in matrix form. The upper and lower bounds $\bar{\mathbf{x}}_{t+T}$ and $\underline{\mathbf{x}}_{t+T}$ together define an axis-aligned hyperrectangle that contains the reachable set $\mathcal{R}_T(\mathcal{X}_t)$.

This containment constraint $\underline{\mathbf{x}}_{t+T} \leq \mathbf{x}_{t+T} \leq \bar{\mathbf{x}}_{t+T}$ can be written as a system of linear inequalities in matrix form:

$$\begin{bmatrix} I \\ -I \end{bmatrix} \mathbf{x}_{t+T} \leq \begin{bmatrix} \bar{\mathbf{x}}_{t+T} \\ -\underline{\mathbf{x}}_{t+T} \end{bmatrix} \quad (27)$$

This elegant representation combines both the upper and lower bounds into a single inequality. The upper part expresses $\mathbf{x}_{t+T} \leq \bar{\mathbf{x}}_{t+T}$ (the upper bounds from our maximization problems), while the lower part expresses $-\mathbf{x}_{t+T} \leq -\underline{\mathbf{x}}_{t+T}$, which is equivalent to $\mathbf{x}_{t+T} \geq \underline{\mathbf{x}}_{t+T}$ (the lower bounds from our minimization problems).

These matrix bounds can be used directly in optimization or verification algorithms, providing a concise way to represent the hyperrectangle that over-approximates the reachable set. In fact, many neural network verification software implementations, such as auto-LiRPA¹ and JAX-Verify², accept a matrix specification, allowing you to plug $\begin{bmatrix} I \\ -I \end{bmatrix}$ directly into these tools and get the complete hyperrectangle bounds in a single function call.

8 More General Bounds

If we choose \mathbf{c} vectors to point in different directions (not just axis-aligned), we can get more general polytope bounds. Let's denote these direction vectors as $\mathbf{c}_1, \mathbf{c}_2, \dots, \mathbf{c}_m$. For each direction \mathbf{c}_i , we solve an optimization problem:

$$d_i = \max_{\mathbf{x}_t \in \mathcal{X}_t} \mathbf{c}_i^T \mathbf{x}_{t+T} \quad \text{s.t.} \quad \mathbf{x}_{t+T} = f^T(\mathbf{x}_t, \pi) \quad (28)$$

We can then stack the vectors \mathbf{c}_i to form a matrix $C = [\mathbf{c}_1^T; \mathbf{c}_2^T; \dots; \mathbf{c}_m^T]$, and collect the optimization results to form a vector $\mathbf{d} = [d_1; d_2; \dots; d_m]$. This gives us the polytope bounds:

$$\overline{\mathcal{R}}_T(\mathcal{X}_t) = \{\mathbf{x}_{t+T} \mid C\mathbf{x}_{t+T} \leq \mathbf{d}\} \supseteq \mathcal{R}_T(\mathcal{X}_t) \quad (29)$$




Figure 5: Illustration of general polytope bounds (red dashed octagon) compared to hyperrectangle bounds (faded red dashed rectangle). By using different objective vectors $\mathbf{c}_1, \mathbf{c}_2, \dots, \mathbf{c}_7$ that are normal to the polytope faces, we can obtain tighter over-approximations of the reachable set $\mathcal{R}_1(\mathcal{X}_t)$ (blue). Each face of the polytope corresponds to a halfspace constraint of the form $\mathbf{c}_i^T \mathbf{x} \leq d_i$, where the \mathbf{c}_i vectors point in the outward normal direction.

Using more general polytope bounds with carefully chosen direction vectors \mathbf{c} allows for tighter approximations of the reachable set compared to axis-aligned hyperrectangles. Each face of the polytope corresponds to a halfspace constraint $\mathbf{c}_i^T \mathbf{x} \leq d_i$, where \mathbf{c}_i is the outward normal vector to that face. The resulting polytope can more closely conform to the shape of the true reachable set, potentially reducing conservatism in verification tasks.

¹https://github.com/KaidiXu/auto_LiRPA

²<https://github.com/google-deepmind/jax-verify>

References

- [1] Katz, G., Barrett, C., Dill, D.L., Julian, K., & Kochenderfer, M.J. (2017). *Reluplex: An Efficient SMT Solver for Verifying Deep Neural Networks*. In Computer Aided Verification. CAV 2017. Lecture Notes in Computer Science, vol 10426, pp. 97–117.